

ECS 150 - Synchronization

Prof. Joël Porquet-Lupine

UC Davis - FQ22



Threads (recap)

Memory sharing

- Private processor registers
- Private stack
- Shared *global* memory

Type of sharing

Independent

- Threads work on distinct areas of shared data

```
int a[N], b[N], c[N], d[N], e[N];

void do_mult(int p, int m)
{
    for (int i = p; i < m; i++)
        a[i] += b[i] * c[i] + d[i] * e[i];
}

int main(void)
{
    ...
    thread_create(do_mult, 0, N/2);
    thread_create(do_mult, N/2, N);
    ...
}
```

Cooperating

- Threads work on same areas of shared data

```
int a[N], b[N], c[N], d[N], e[N];

void do_mult_bc(void)
{
    for (int i = 0; i < n; i++)
        a[i] += b[i] * c[i];
}

void do_mult_de(void)
{
    for (int i = 0; i < n; i++)
        a[i] += d[i] * e[i];
}

int main(void)
{
    ...
    thread_create(do_mult_bc);
    thread_create(do_mult_de);
    ...
}
```

Concurrency issues

Example

```
int x = 0;

void thread_a(void)
{
    x = x + 1;
}

void thread_b(void)
{
    x = x + 2;
}

int main(void)
{
    thread_create(thread_a);
    thread_create(thread_b);
    thread_join(thread_a);
    thread_join(thread_b);
    printf("%d\n", x);
    return 0;
}
```

Execution

- Typical output:

```
$ ./a.out
3
```

- Also possible (yet probably very rare)...


```
$ ./a.out
2
$ ./a.out
1
```

Concurrency issues

Indeterministic interleavings

Thread scheduling

- Indeterministic scheduling
- Sequential execution, concurrent execution, parallelism, etc.



Instruction reordering

- Compiler instruction reordering
- Hardware instruction reordering

```
void thread1(void) {  
    ...  
    p = init();  
    init = true;  
    ...  
}  
  
void thread2(void) {  
    ...  
    while (init == false);  
    q = docomputation(p);  
    ...  
}
```

Multi-word operations

- Multi-word operations are not atomic

```
// Assuming 32-bit CPU  
uint64_t num = 2;  
...  
num = num + 1;
```

Concurrency issues

Race conditions

Definition

- **Race condition** when the result of a concurrent program depends on the order of operations between threads.

```
void thread_a(void)
{
    x = x + 1;
}
```

```
void thread_b(void)
{
    x = x + 2;
}
```

Difficulties

- Number of possible "interleavings" can be huge
- Some interleavings are good
 - Vast majority of them usually are
- Some interleavings are bad
 - They may even be extremely rare...

Solution?

- Synchronization!

Too Much Milk!

Example

| Roommate 1 | Roommate 2 |
|-----------------------------|-----------------------------|
| Arrive home | |
| Look in fridge, out of milk | |
| Leave for store | |
| | Arrive home |
| Arrive at store | Look in fridge, out of milk |
| Buy milk | Leave for store |
| Arrive home, put milk away | Buy milk |
| | Arrive home, put milk away |
| | Oh, no! |

Required correctness properties

- **Safety:** at most one person buys milk at a time
- **Liveness:** someone buys milk if needed

Too Much Milk!

1. Leaving a note

Roommate 1 (thread 1)

```
if (note == 0) {  
    if (milk == 0) {  
        note = 1;  
        milk++;  
        note = 0;  
    }  
}
```

Roommate 2 (thread 2)

```
if (note == 0) {  
    if (milk == 0) {  
        note = 1;  
        milk++;  
        note = 0;  
    }  
}
```

Safety and liveness

- Not safe if threads are descheduled right after the two tests
 - They would both put a note and get milk, resulting in two bottles of milk!

Too Much Milk!

2. Using two notes

Roommate 1 (thread 1)

```
note1 = 1;
if (note2 == 0) {
    if (milk == 0) {
        milk++;
    }
}
note1 = 0;
```

Roommate 2 (thread 2)

```
note2 = 1;
if (note1 == 0) {
    if (milk == 0) {
        milk++;
    }
}
note2 = 0;
```

Safety and liveness

- Not live if threads are descheduled right after setting their personal notes
 - They both think the other is on their way getting milk but no one eventually does

Too Much Milk!

3. Using asymmetric notes

Roommate 1 (thread 1)

```
note1 = 1;
while (note2 == 1)
    ;
if (milk == 0) {
    milk++;
}
note1 = 0;
```

Roommate 2 (thread 2)

```
note2 = 1;
if (note1 == 0) {
    if (milk == 0) {
        milk++;
    }
}
note2 = 0;
```

Safety and liveness

- Yes, safe and live!

Too Much Milk!

3. Using asymmetric notes: explained

Roommate 1 (thread 1)

```
note1 = 1;
while (note2 == 1)
    ;
if (milk == 0) {
    milk++;
}
note1 = 0;
```

Roommate 2 (thread 2)

```
note2 = 1;
if (note1 == 0) {
    if (milk == 0) {
        milk++;
    }
}
note2 = 0;
```

Issues

1. Way too over-engineered!
2. Asymmetrical, non-scalable
3. Involves *busy-waiting*

Peterson's algorithm

- Share a single resource using only shared memory
- Symmetric and scalable
- But still quite complex...
- [More here](#)

Critical section

Definition

- Piece of code where the shared resource is accessed
- Needs to be properly protected to avoid race conditions
- Cannot be executed by more than one thread at a time

Thread 1

```
...  
CS_enter();  
    Critical section  
CS_exit();  
...
```

Thread 2

```
...  
CS_enter();  
    Critical section  
CS_exit();  
...
```

Correctness properties

1. **Safety**
2. **Liveness**
3. Bounded waiting
4. Failure atomicity

Mutual exclusion

- Property of concurrency control
- Requirement that only one thread can enter critical section at a time
- Active thread excludes its peers

Critical section

Formalizing "Too Much Milk!"

- Shared variable
- Safety property
- Liveness property

Roommate 1 (thread 1)

```
note1 = 1;           /*  
while (note2 == 1)    * CS_enter()  
    ;                */  
if (milk == 0) {      /*  
    milk++;           * CS  
}                     */  
note1 = 0;           /* CS_exit() */
```

Roommate 2 (thread 2)

```
note2 = 1;           /* CS_enter()  
if (note1 == 0) {     */  
    if (milk == 0) {   /*  
        milk++;       * CS  
    }                 */  
}                     */  
note2 = 0;           /* CS_exit() */
```

ECS 150 - Synchronization

Prof. Joël Porquet-Lupine

UC Davis - FQ22



Recap


Race condition

- Output of concurrent program depends on order of operations between threads

```
void thread_a(void)      void thread_b(void)
{
    x = x + 1;
}
{
    x = x + 2;
}
```

```
$ ./a.out
3
$ ./a.out # Very rare but possible...
2
$ ./a.out # Very rare but possible...
1
```

- Indeterministic concurrent execution
 - Thread scheduling
 - Instruction reordering
 - Multi-word operations



Critical section

```
void thread_1(void) {
    note1 = 1;           // v
    while (note2 == 1)    // CS_enter()
        ;                // ^
    if (milk == 0) {      // v
        milk++;           // CS
    }                     // ^
    note1 = 0;           // CS_exit()
}

void thread_2(void) {
    note2 = 1;           // CS_enter()
    if (note1 == 0) {     // ^
        if (milk == 0) {  // v
            milk++;       // CS
        }                 // ^
    }
    note2 = 0;           // CS_exit()
}
```

- Shared variable
- Safety property
- Liveness property
- Mutual exclusion

Locks

Definition

- A lock is a *synchronization* variable that provides *mutual exclusion*
- Two states: *locked* and *free* (initial state is generally *free*)



API

- `lock()` or `acquire()`
 - Wait until lock is free, then grab it
- `unlock()` or `release()`
 - Unlock, and allow one of the threads waiting in *acquire* to proceed

```
int milk;

int main(void)
{
    thread_create(roommate_fcn);
    thread_create(roommate_fcn);
    ...
}
```

```
void roommate_fcn(void)
{
    ...
    lock();

    /* Critical section */
    if (!milk)
        milk++;

    unlock();
    ...
}
```

Locks

Simple uniprocessor implementation

- Race conditions are coming from indeterministic scheduling
 - Breaks atomicity of instruction sequence
 - Caused by preemption (i.e. timer interrupt)
- Solution: disable the interrupts!

```
void lock(void)
{
    disable_interrupts();
}
```

```
void unlock(void)
{
    enable_interrupts();
}
```

Issues


- Only works on uniprocessor systems
- Dangerous to have unpreemptable code
- Cannot be used by user applications

Multiprocessor spinlocks

Hardware support

- *Test-and-set* hardware primitive to provide mutual exclusion
 - a.k.a *Read-modify-write* operation
- Typically relies on a multi-cycle bus operation that **atomically** reads and updates a memory location
 - Multiprocessor support

```
/* Equivalent of a test&set hardware
instruction in software */
ATOMIC int test_and_set(int *mem)
{
    int oldval = *mem;
    *mem = 1;
    return oldval;
}
```



Lock implementation

```
void spinlock_lock(int *lock)
{
    while (test_and_set(lock) == 1);
}
```

```
void spinlock_unlock(int *lock)
{
    *lock = 0;
}
```

Multiprocessor spinlocks

Revisiting "Too Much Milk!"

```
int lock = 0;
```

Thread 1

```
spinlock_lock(&lock);  
if (milk == 0) {  
    milk++;  
}  
spinlock_unlock(&lock);
```

Thread 2

```
spinlock_lock(&lock);  
if (milk == 0) {  
    milk++;  
}  
spinlock_unlock(&lock);
```

Thread 3

```
spinlock_lock(&lock);  
if (milk == 0) {  
    milk++;  
}  
spinlock_unlock(&lock);
```

Thread 4

```
spinlock_lock(&lock);  
if (milk == 0) {  
    milk++;  
}  
spinlock_unlock(&lock);
```

Multiprocessor spinlocks

Issue

- Busy-waiting wastes cpu cycles
 - Only to reduce latency

```
void spinlock_lock(int *lock)
{
    while (test_and_set(lock) == 1);
}
```

Solution

"Cheap" busy-waiting

- Yield/sleep when unable to get the lock, instead of looping

```
void lock(int *lock)
{
    while (test_and_set(lock) == 1)
        thread_yield(); //or, sleep(N);
}
```

Better primitives

- *Block* waiting threads until they can proceed

```
void lock(int *lock)
{
    while (test_and_set(lock) == 1)
        thread_block(lock);
}
```

Cons

- Yielding still wastes cpu cycles
- Sleeping impacts latency as well


Examples

- Semaphores
- Mutexes (equivalent to binary semaphore with the notion of ownership)

Semaphores

Definition

- Invented by Dijkstra in the 60's
- A semaphore is a generalized lock
 - Used for different types of synchronization (including mutual exclusion)
 - Keeps track an arbitrary resource count
 - Queue of threads waiting to access resource



One resource to share



Multiple resources to share

Semaphores

API

- Initial count value (*but **not** a maximum value*)

```
sem = sem_create(count);
```

- `down()` or `P()`
 - Decrement by one, or block if already 0
- `up()` or `V()`
 - Increment by one, and wake up one of the waiting threads if any

Possible implementation:

```
void sem_down(sem)
{
    spinlock_lock(sem->lock);
    while (sem->count == 0) {
        /* Block self */
        ...
    }
    sem->count -= 1;
    spinlock_unlock(sem->lock);
}
```

```
void sem_up(sem)
{
    spinlock_lock(sem->lock);
    sem->count += 1;
    /* Wake up first in line */
    /* (if any) */
    ...
    spinlock_unlock(sem->lock);
}
```

Semaphores

Binary semaphore

- Semaphore which count value is either 0 or 1
- Can be used similarly as a lock
 - But no busy waiting, waiting thread are blocked until they can get the lock
- Guarantees mutually exclusive access to a shared resource
- Initial value is generally 1 (ie *free*)

Example

```
sem = sem_create(1);
```

Thread 1

```
...  
down(sem);  
    Critical section  
up(sem);  
...
```

Thread 2

```
...  
down(sem);  
    Critical section  
up(sem);  
...
```

Semaphores

Counted semaphore

- Semaphore which count value can be any *positive* integer
 - Represents a resource with many "units" available
- Initial count is often the number of initial resources (if any)
- Allows a thread to continue as long as enough resources are available
- Used for synchronization

Example

```
sem_packet = sem_create(0);
```

Thread 1

```
while (1) {  
    x = get_network_packet();  
    enqueue(packetq, x);  
    up(sem_packet);  
}
```

Thread 2

```
while (1) {  
    down(sem_packet);  
    x = dequeue(packetq);  
    process_contents(x);  
}
```

ECS 150 - Synchronization

Prof. Joël Porquet-Lupine

UC Davis - FQ22



Recap

Locks

```
void thread(void) {  
    lock();  
    /* Critical section */  
    ...  
    unlock();  
}
```

Atomic spinlocks

- Based on atomic *test-and-set* instruction
- Compatible with multiprocessor systems
- Accessible to user processes

```
void spinlock_lock(int *lock) {  
    while (test_and_set(lock) == 1);  
}
```

```
void spinlock_unlock(int *lock) {  
    *lock = 0;  
}
```

- But based on busy-waiting

Semaphores

- Internal count
- `down()` decrements count by one, or blocks if count is 0
- `up()` increments count by one, and wakes up first blocked thread if any

Binary semaphore

```
sem = sem_create(1);  
void thread(void) {  
    sem_down(sem);  
    /* Critical section */  
    ...  
    sem_up(sem);  
}
```


Counted semaphore

```
sem = sem_create(0);  
void thread1(void) {  
    x = get_packet();  
    enqueue(q, x);  
    up(sem);  
}  
void thread2(void) {  
    down(sem);  
    x = dequeue(q);  
    process_packet(x);  
}
```

Producer-consumer problem

Definition

Two or more threads communicate through a circular data buffer: some threads *produce* data that others *consume*.



- Bounded buffer of size N
- Producers write data to buffer
 - Write at `in` and moves rightwards
 - Don't write more than the amount of available space
- Consumers read data from buffer
 - Read at `out` and moves rightwards
 - Don't consume if there is no data
- Allow for multiple producers and consumers

Producer-consumer problem

Solution 1: no protection

```
int buf[N], in, out;
```

```
void produce(int item)
{
    buf[in] = item;
    in = (in + 1) % N;
}

void thread_prod(void)
{
    while (1) {
        ...
        produce(x);
        ...
    }
}
```

```
int consume(void)
{
    int item = buf[out];
    out = (out + 1) % N;
    return item;
}

void thread_cons(void)
{
    while(1) {
        ...
        x = consume();
        ...
    }
}
```

Issues

- Unprotected shared state
 - Race conditions on all shared variables
- No synchronization between consumers and producers

Producer-consumer problem

Solution 2: Lock semaphores

- Add protection of share state
 - Mutual exclusion around critical sections
 - Guarantees one producer and one consumer at a time

```
int buf[N], in, out;  
sem_t lock_prod = sem_create(1), lock_cons = sem_create(1);
```

```
void produce(int item)  
{  
    sem_down(lock_prod);  
    buf[in] = item;  
    in = (in + 1) % N;  
    sem_up(lock_prod);  
}
```

```
int consume(void)  
{  
    sem_down(lock_cons);  
    int item = buf[out];  
    out = (out + 1) % N;  
    sem_up(lock_cons);  
    return item  
}
```

Producer-consumer problem

Solution 3: Communication semaphores

- Add synchronization between producers and consumers
 - Producers wait if buffer is full
 - Consumers wait if buffer is empty

```
int buf[N], in, out;  
sem_t lock_prod = sem_create(1), lock_cons = sem_create(1);  
sem_t empty = sem_create(N), full = sem_create(0);
```

```
void produce(int item)  
{  
    sem_down(empty); //need empty spot  
    sem_down(lock_prod);  
    buf[in] = item;  
    in = (in + 1) % N;  
    sem_up(lock_prod);  
    sem_up(full); //new item avail  
}
```

```
int consume(void)  
{  
    sem_down(full); //need new item  
    sem_down(lock_cons);  
    int item = buf[out];  
    out = (out + 1) % N;  
    sem_up(lock_cons);  
    sem_up(empty); //empty slot avail  
    return item  
}
```




Readers-writers problem

Definition

- Multiple threads access the same shared resource, but differently
 - *Many* threads only *read* from it
 - *Few* threads only *write* to it
- Readers vs writers
 - Multiple *concurrent* readers at a time
 - Or single writer at a time

Examples

- Airline ticket reservation
- File manipulation



Readers-writers problem

Solution 1: Protect resource

```
sem_t rw_lock = sem_create(1);
```

```
void writer(void)
{
    sem_down(rw_lock);
    ...
    /* perform write */
    ...
    sem_up(rw_lock);
}
```

```
int reader(void)
{
    sem_down(rw_lock);
    ...
    /* perform read */
    ...
    sem_up(rw_lock);
}
```

Analysis

- Mutual exclusion between readers and writers: *Yes*
- Only one writer can access the critical section: *Yes*
- Multiple readers can access the critical section at the same time: *No!*

Readers-writers problem

Solution 2: Enable multiple readers

```
int rcount = 0;
sem_t rw_lock = sem_create(1);
```

```
void writer(void)
{
    sem_down(rw_lock);
    ...
    /* perform write */
    ...
    sem_up(rw_lock);
}
```

```
int reader(void)
{
    rcount++;
    if (rcount == 1)
        sem_down(rw_lock);
    ...
    /* perform read */
    ...
    rcount--;
    if (rcount == 0)
        sem_up(rw_lock);
}
```

Issue

- Race condition between readers on variable `rcount`!

Readers-writers problem

Solution 3: Protect multiple readers

```
int rcount = 0;
sem_t rw_lock = sem_create(1), count_lock = sem_create(1);
```

```
void writer(void)
{
    sem_down(rw_lock);
    ...
    /* perform write */
    ...
    sem_up(rw_lock);
}
```

Analysis

- Correct solution
- But suffers from potential starvation of writers

```
int reader(void)
{
    sem_down(count_lock);
    rcount++;
    if (rcount == 1)
        sem_down(rw_lock);
    sem_up(count_lock);

    ...
    /* perform read */
    ...
    sem_down(count_lock);
    rcount--;
    if (rcount == 0)
        sem_up(rw_lock);
    sem_up(count_lock);
}
```

Semaphores

Concluding notes

- *Semaphores considered harmful* (Dijkstra, 1968)
 - Simple algorithms can require more than one semaphore
 - Increase of complexity to manage them all
 - Semaphores are low-level primitives
 - Easy to make programming mistakes (e.g. `down ()` followed by `down ()`)
 - Programmer must keep track of the order of all semaphores operations
 - Avoid deadlocks
 - Semaphores are used for both mutual exclusion and synchronization between threads
 - Difficult to determine which meaning a given semaphore has
- Need for another abstraction
 - Clear distinction between mutual exclusion and synchronization aspects
 - Concept of *monitor* developed in early 70's


Synchronization barriers

Concept

- Enables multiple threads to wait until all threads have reached a particular point of execution before being able to proceed further

```
void main(void)
{
    barrier_t b = barrier_create(3);
    thread_create(thread_func, b);
    thread_create(thread_func, b);
    thread_create(thread_func, b);
}

void thread_func(barrier_t b)
{
    while (/* condition */) {
        /* ... do some computation ... */
        ...
        /* Wait for the other threads */
        barrier_wait(b);
    }
}
```



Implementation

- Using semaphores
- Using condition variables and the `broadcast()` feature

Synchronization: the big picture

Concurrent applications

| | | |
|----------------------------------|---------------------|--------------------------------------------------------------|
| <i>Shared Objects</i> | Bounded buffer | Barrier |
| <i>Synchronization Variables</i> | Semaphores | <div><i>Monitor</i> Locks Condition variables</div> |
| <i>Atomic Instructions</i> | Interrupt disabling | Test-and-set instructions |
| <i>Hardware</i> | Hardware interrupts | Multiprocessors |

Best practices

- [Basic Threads Programming: Standards and Strategy](#).
- [The 12th Commandments of Synchronization](#) (Cornell University, 2011)

Annex: Monitors

Definition

A monitor is an abstract data type that controls access to shared resources.

It is comprised of:

- Shared private data
 - The resource to protect
- Procedures that operate on the data
 - Gateway to the resource
- Synchronization primitives
 - Enforces the sequence of access
 - *Condition variables*

Queue example

```
monitor queue {  
    integer size;  
    condition fullCV, emptyCV;  
  
    procedure enqueue() {  
        ...  
    }  
  
    procedure dequeue() {  
        ...  
    }  
}
```

Annex: Monitors

Mutual exclusion

A monitor defines an internal lock for providing mutual exclusion.

- Ensures that only a single thread is active in the monitor at any point
- Also provides mutual exclusion for shared data

```
monitor queue {  
    integer size;  
    condition fullCV, emptyCV;  
    lock l;  
  
    procedure enqueue() {  
        l.lock();  
        ...  
        l.unlock();  
    }  
  
    procedure dequeue() {  
        l.lock();  
        ...  
        l.unlock();  
    }  
}
```

Annex: Monitors

Condition variable

Definition

A monitor defines zero or more condition variables for providing synchronization.

- Enables threads to wait efficiently for changes to shared data, while protected by a lock
- Each condition variable is essentially a queue of waiting threads (but stateless compared to semaphores)

API

- `wait()`: Put calling thread in wait queue
- `signal()`: Wake up first thread in wait queue, if any

Annex: Monitors

Condition variable

Example

```
monitor queue {  
    integer size;  
    condition fullCV, emptyCV;  
    lock l;  
  
    procedure enqueue(item) {  
        l.lock();  
  
        queue.add(item);  
  
        size++;  
  
        fullCV.signal();  
  
        l.unlock();  
    }  
}
```

```
    procedure dequeue(&item) {  
        l.lock();  
  
        while (size == 0)  
            fullCV.wait(l);  
  
        size--;  
  
        item = queue.remove();  
  
        l.unlock();  
    }  
}
```

Rule: a thread **must** hold the monitor lock when doing condition variable operations!

Annex: Monitors

Condition variable

API

- `wait(lock l)`
 - Release lock, move thread to waiting queue, suspend thread (atomically)
 - When thread wakes up, it re-acquires the lock before returning from wait
 - Condition predicate must be rechecked because of possible spurious wake-up
- `signal()`
 - Wake up the first waiting thread, if any
 - What happens if no waiting thread?
- `broadcast()`
 - Wake up all waiting threads, if any
 - What happens if no waiting threads?

Annex: Monitors

Condition variable

signal() semantic

```
procedure enqueue(item) {  
    ...  
    size++;  
    fullCV.signal();  
    ...  
}
```

Hoare

- Original behavior (often presented in textbooks)
- `signal()` immediately switches from the signaler to a waiting thread
- The condition that the waiter was anticipating is guaranteed to hold when waiter executes
 - I.e., could use `if` once to check the wait condition

```
procedure dequeue(&item) {  
    ...  
    while (size == 0)  
        fullCV.wait(1);  
    size--;  
    ...  
}
```

Mesa

- Most existing implementations
- `signal()` places the waiter on the ready queue, but signaler continues inside the monitor
- Condition is not necessarily true when waiter runs again
 - Returning from `wait()` is only a hint that something has changed
 - Need to check condition again

Annex: Monitors

Producer-consumer problem

```
monitor prod_cons {
    lock l;
    condition fullCV, emptyCV;
    item buffer[N];
    integer empty = N, full = 0;
    integer in = 0, out = 0;

    procedure produce(item) {
        l.lock();

        while(empty == 0)
            emptyCV.wait(l);
        empty--;

        buffer[in] = item;
        in = (in + 1) % N;

        full++;
        fullCV.signal();

        l.unlock();
    }
}
```

```
procedure consume(&item) {
    l.lock();

    while (full == 0)
        fullCV.wait(l);
    full--;

    item = buffer[out];
    out = (out + 1) % N;

    empty++;
    emptyCV.signal();

    l.unlock();
}
}
```

Annex: Monitors

Programming language inclusion

- First introduced as a programming language construct (Mesa, Java)
 - Embedded in the language itself
 - Added by the compiler and enforced at runtime

```
class Queue extends Object {  
    synchronized int read() {  
        while(no data) {  
            this.wait();  
        }  
        ...  
    }  
    synchronized void write(int i) {  
        ...  
        this.notify();  
    }  
}
```

- More largely, defines a programming convention
 - Can be recreated in most languages (e.g. C/C++ with pthread library)
 - Condition variables can be implemented [using semaphores](#)

Annex: Monitors

Vs semaphores

Both

- Provide mutual exclusion and synchronization
- Have `signal()/up()` and `wait()/down()`
- Support a queue of threads that are waiting to access a critical section
- No busy waiting

Semaphores

- Semaphores are essentially generalized locks
- Binary and counting semaphores (stateful)
- Used for mutual exclusion and synchronization

Monitors

- Consist of a lock and zero or more condition variables
- Encapsulate shared data
- Use lock for mutual exclusion
- Use condition variable(s) for synchronization (stateless)
- Wrap operations on shared data with lock
- Condition variables release lock temporarily