

ECS 150 - Virtual Memory: Demand Paging

Prof. Joël Porquet-Lupine

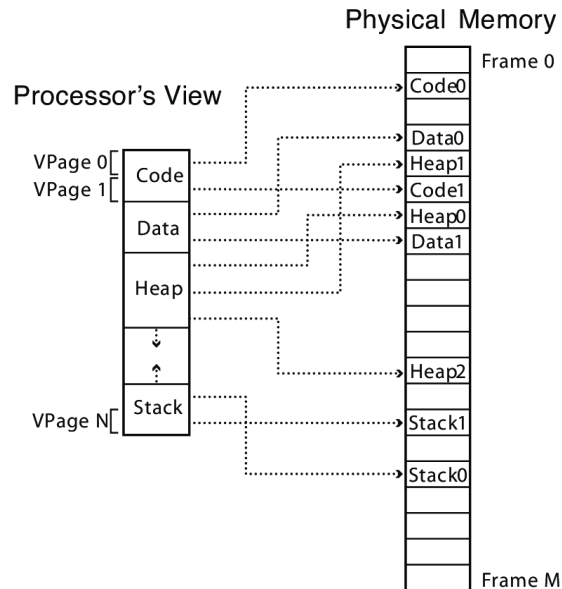
UC Davis - FQ22



Introduction

Address translation

- Distinction between virtual address space and physical memory
 - Mapping at page level providing relocation and protection
- Virtual address space convenient for programs and processes
 - Segments contiguously allocated
 - At addresses allocated by compiler
- Physical memory easy to allocate
 - Any available page frame can be mapped to any process virtual page



Introduction

Issue

- Assumption that a running process has its entire address space loaded
 - Code, data, heap, and stack segments
- Slow, especially for big processes
 - Code and data need to be loaded from disk
- Wasteful
 - Processes don't use all of their memory all the time
 - *Working set* is usually small and evolves slowly

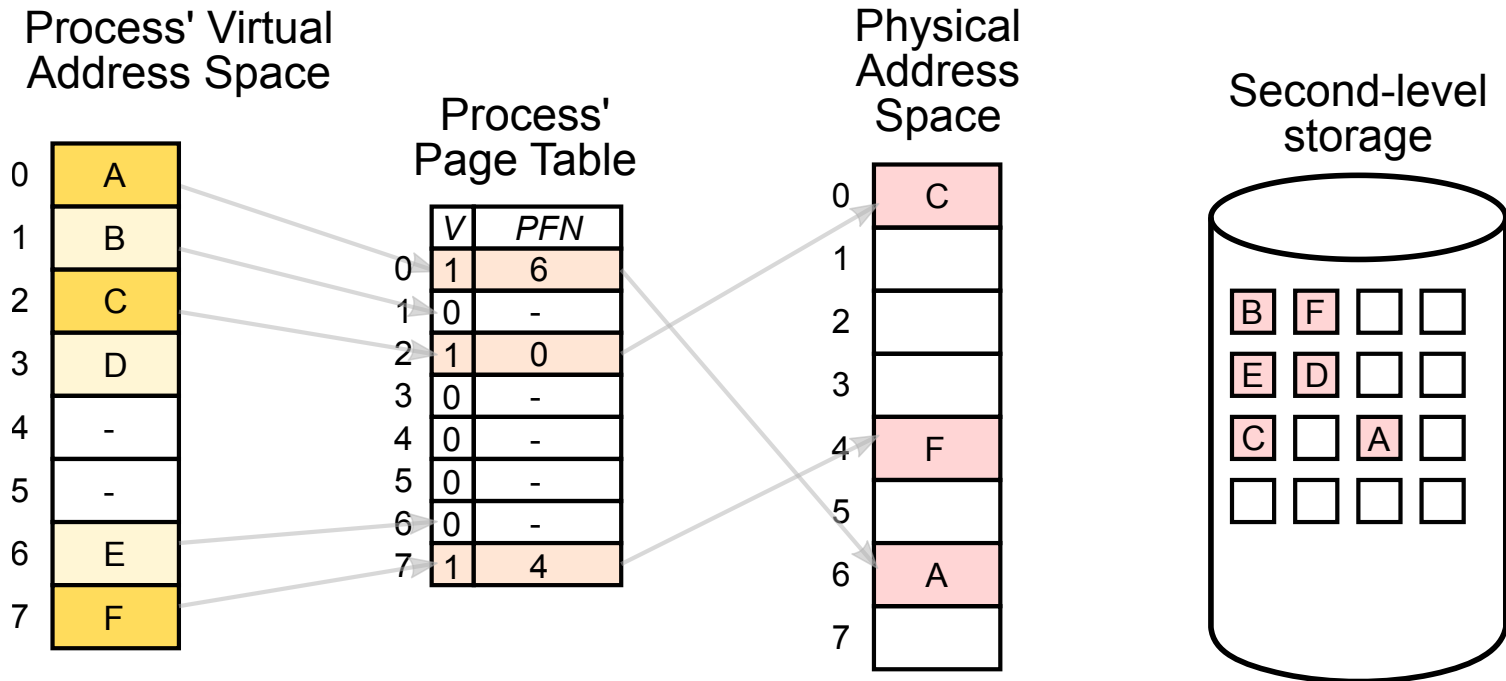
Idea

- Load pages as needed during process execution
 - Only bring in pages actually used
 - Only keep frequently-used pages in memory
- Illusion of (nearly) infinite memory, available to every process
- Multiplex virtual pages onto a limited amount of physical page frames

Demand paging

Page mapping

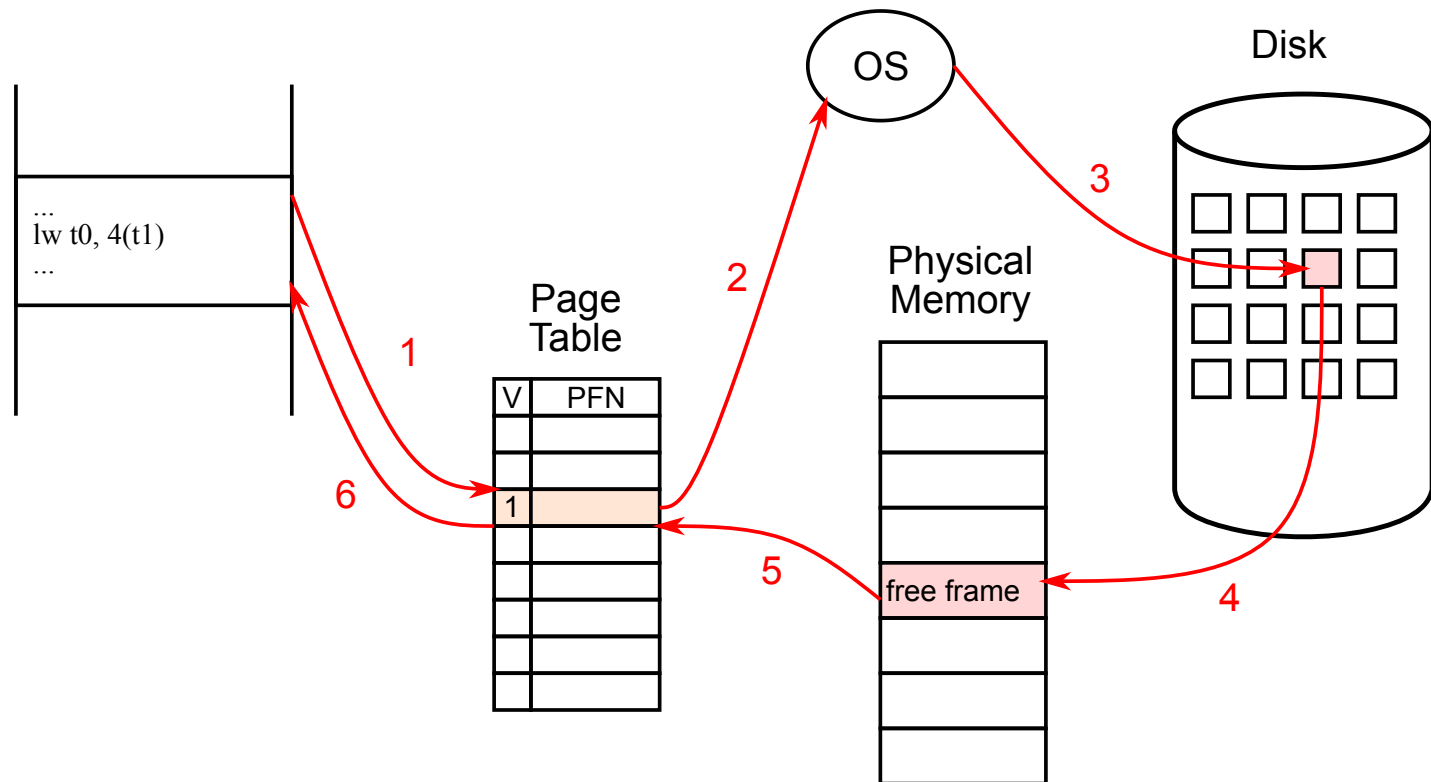
- *Resident* page: mapped in physical memory, valid PTE
- *Non-resident* page: located on disk, invalid PTE
 - Page-fault triggered if accessed
- *Unused* page



Demand paging

Page fault (overview)

- (1) Memory access, TLB access
- (2) Page invalid, trap to kernel
- (3) Locate page on disk
- (4) Swap in page in free frame
- (5) Mark page as valid
- (6) Resume process



Demand paging

Page fault (details)

- (1) Memory access from running process
 - Lookup virtual memory address in TLB
 - If TLB miss, page table walk
- (2) Page is invalid
 - Page fault, trap to kernel
- (3) Locate page on disk
 - Directly from executable if page of code
 - Reuse PTE to store the block number in swap space
- (4) Swap in page in free frame
 - Allocate page frame
 - Evict page if needed
- (4) Cont.
 - Initiate disk block read into page frame
 - Elect another process to run...
 - Disk interrupt when DMA complete
- (5) Mark page as valid
 - Purge TLB for this page
- (6) Resume process
 - At faulting instruction
 - TLB miss
 - Page table walk to fetch latest translation
 - Execute instruction and access page

Demand paging

Page frame allocation

- If memory isn't full, allocating free page frame is straight-forward
- Otherwise, need a specific strategy
 - Select old page to evict
 - *Page replacement* algorithm (e.g., FIFO, LRU, etc.)
 - Unmap old page from processes
 - Page frame can be shared among several processes
 - Find all page table entries that refer to old page
 - Set page table entries to invalid
 - Purge corresponding TLB entries
 - Write changes on page back to disk, if necessary
 - Need to detect **page modifications**

```
fd = open("/path/to/file");
char *address = mmap(0, len, PROT_WRITE,
                     MAP_SHARED, fd, 0)
// Write
address[some_offset] = 'a';
...
```

Will need to write back (dirty) page to disk

```
fd = open("/path/to/file");
char *address = mmap(0, len, PROT_READ,
                     MAP_SHARED, fd, 0)
// Only read
printf("%x\n", address[some_offset]);
...
```

Can simply discard page when done

Demand paging

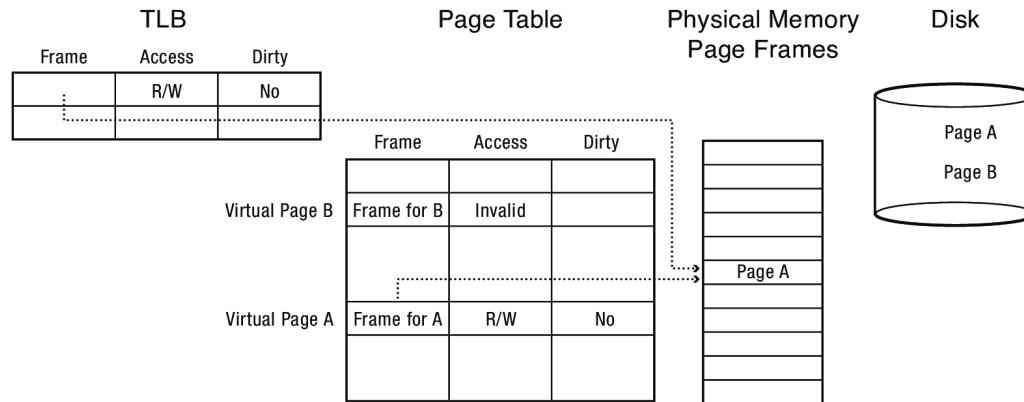
PTE bookkeeping

- On most processor architectures, every PTE has some bookkeeping support
 - **Modified bit**
 - Set by hardware in TLB entry on store instruction
 - Usually updated to PTE upon TLB eviction
 - **Use bit** (also known as *Reference bit*)
 - Set by hardware in PTE upon TLB miss
- Bookkeeping bits can be reset by OS
 - When changes to page are purged to disk
 - To track whether page has recently been used

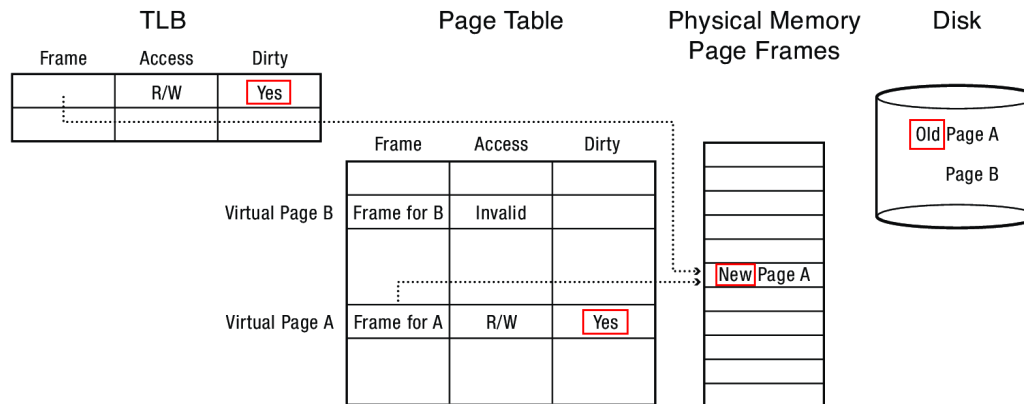
Demand paging

Tracking page modifications

Before to clean page:



After writing to page for the first time:



Page replacement

Strategies

- When memory is full, need to select a *victim frame* to evict
- Algorithms
 - Random
 - Zero cost for bookkeeping
 - Not best, not worst, just unpredictable!
 - FIFO
 - MIN
 - LRU

Page replacement

FIFO

- Replace the page that has been loaded the longest time

Example

- 4 page frames available
- Access sequence: A B C D A B E A B C D E

Reference	A	B	C	D	A	B	E	A	B	C	D	E
Frame #1	A				+		E				D	
Frame #2		B				+		A				E
Frame #3			C						B			
Frame #4				D						C		

- Result: **10 page faults**

Pros

- Simple implementation

Cons

- May replace the heavily-used pages
- Suffers from *Belady's anomaly*

Page replacement

Belady's anomaly

- FIFO algorithm with 4 page frames causes 10 page faults

Reference	A	B	C	D	A	B	E	A	B	C	D	E
Frame #1	A				+		E				D	
Frame #2		B				+		A				E
Frame #3			C						B			
Frame #4				D						C		

- Now, same sequence but with only **3** page frames available...

Reference	A	B	C	D	A	B	E	A	B	C	D	E
Frame #1	A			D			E					+
Frame #2		B			A			+		C		
Frame #3			C			B			+		D	

- 9 page faults!**
 - Less page faults although less page frames are available...!?

Page replacement

MIN (aka optimal)

- Replace the page that will not be used for the longest time in the future

Example

Reference	A	B	C	D	A	B	E	A	B	C	D	E
Frame #1	A				+			+			D	
Frame #2		B				+			+			
Frame #3			C							+		
Frame #4				D			E					+

Pros

- Only 6 faults, optimal!

Cons

- Impossible to implement, only gives the ideal lower bound

Page replacement

LRU

- Replace the page that has not been used for the longest time in the past

Example

Reference	A	B	C	D	A	B	E	A	B	C	D	E
Frame #1	A				+			+				E
Frame #2		B				+			+			
Frame #3			C				E				D	
Frame #4				D						C		

Pros

- Good approximation of MIN

Cons

- Difficult to accurately implement

Page replacement

Implementing LRU

Issue

- In software, use a linked list
 - Every hit moves the page to the front of the list
 - Evict from the back of the list
 - But you can't trap back to kernel for each page access...
- In hardware
 - Impossible to manage a variable-size linked list...

Solution

- Approximate LRU
- Take advantage of the *use bit*!


Page replacement

Clock

- Periodically, sweep through all the pages
- If page is unused, reclaim
- If page is used, mark as unused

Second-chance

- Significant cost to reclaim *dirty* pages
- Modify clock algorithm to allow dirty pages to survive the first sweep of the clock hand



The end! 🎓
